

Do we need more bikes?

Albert Bargalló I Sales

Michel Misso

Markus Lindgren

William Berger

December 19, 2024

Abstract

This paper explores different machine learning models for predicting whether there will be a low or high demand on shared bicycles in Washington DC. The explored models include Logistic Regression, K-nearest neighbors, Random forest, XGBoost and CatBoost. Hyperparameter tuning was done on all the models and after evaluating them on a test-set with recall as focus metric, random forest was deemed the winner. The random forest model achieved a recall score of 64% and overall accuracy of 88%. The Random Forest model was used to predict on the final test set.

1 Introduction

To reduce dependence on fossil fuel-based transportation, Capital Bikeshare offers Washington DC, a public bicycle-sharing system as a sustainable alternative. However, on certain occasions, the demand for bicycles exceeds the availability, which in turn leads to people choosing less sustainable modes of transportation. Thus, the District Department of Transportation in Washington seeks to fix this problem by providing more bicycles during times of high demand. This project aims to use a machine-learning approach to provide insight in when it would be beneficial to increase the number of bicycles. By creating multiple models for binary classification and evaluating them against each other, the goal is to use the best model to predict if there is high or low demand for bicycles each hour.

2 Data analysis

2.1 Exploring features

Through an initial exploratory analysis of the dataset provided, the structure and types of features available for modeling were examined. This analysis involved identifying the numerical and categorical features.

```

"Numerical": ["temp", "dew", "humidity", "precip", "snow",
              "snow_depth", "windspeed", "cloudcover", "visibility"],
"Categorical": ["hour_of_day", "day_of_week", "month", "holiday",
               "weekday", "summertime"]

```

2.2 Trends

As it can be seen in Figure 1, comparing the demand of bikes on different hours it can be derived that the highest demand occurs on the afternoon (15-19h), and that there is also a higher demand during the morning (8-14h) than at night (20-7h). Analyzing for each month, It seems that there is a considerably higher demand of bikes from March to October than from November to February.




Figure 1: Bike demand plotted against temporal data.

Table 1 highlights a clear difference in high bike demand between weekends and weekdays. The percentage of high bike demand during weekends is 25%, significantly higher compared to 15.14% on weekdays. This indicates that weekends experience a notable increase in bike usage. Table 2 shows that high bike demand during non-holidays (18.03%) is only slightly higher than during holidays (16.98%), indicating that the "holiday" variable had minimal impact on bike demand.

Table 1: High demand by day of week.

Name	Description
Weekends:	25%
Weekdays:	15.14%

Table 2: High demand by holiday.

Name	Description
Non-holiday:	18.03%
Holiday:	16.98%

The histograms of temperature and windspeed in Figure 2 provide insights into the relationship between these variables and bike demand. The windspeed histogram shows a similar distribution pattern for both low and high bike demand, with a larger volume of data corresponding to low bike demand. This indicates that windspeed does not strongly differentiate between the two demand categories, suggesting it may not be a key predictor for high bike demand. In contrast, the temperature histogram reveals a distinct trend. Higher temperatures are clearly associated with increased bike demand, as evidenced by the greater proportion

of high bike demand at elevated temperatures. Conversely, lower temperatures predominantly correspond to low bike demand. This trend underscores the significant role that temperature plays in influencing bike demand, making it a critical feature in predicting demand patterns. These observations reinforce the importance of focusing on variables like temperature that exhibit strong predictive signals, while less impactful variables like windspeed and holiday may contribute less to the model's performance and could be removed.




Figure 2: Bike demand plotted against weather data.

2.3 Feature selection

Based on the data analysis above, a new dataset was created with the most important features. A correlation matrix and plots similar to the figures above were created to help with this process.

First, the column *holiday* was removed based on the minimal information it contained. The column *snow* was also dropped since the entire column only contained 0:s which gives us no additional information on the demand of bikes. Secondly, the *snowdepth* feature was converted to a binary feature *snow*, 1 if there was any snow and 0 otherwise, as it seems to be a significant difference in the demand when there is snow and when there isn't, but not between the amount of snow. *Precip* was also converted to a binary feature *rain*, since the amount rainy days with high bicycle demand was not many. In addition, it was decided to keep the *summertime* feature since there is a clear difference in the demand between summer/winter. Lastly, a new binary feature called *day* was created, being 1 if *hour_of_day* was between 7 and 20h, as it could be useful to represent the difference on demand by *hour_of_day* seen on Figure 1.

Furthermore, the numerical features in the dataset were normalized through *Z-score scaling* to ensure that they were centered around 0 with a standard deviation of 1. The reason for this was to help the model converge quickly and possibly give more accurate predictions.[1] The normalization step is particularly

important for logistic regression and k-NN since they rely on distance measurement to learn. If the feature ranges are very different, the calculations can be dominated by features with larger ranges, leading to biased or suboptimal results.[1] We also decided not to one-hot encode any categorical values since *hour_of_day*, *month*, and *day_of_week* all have a natural ordering. One-hot encoding also increases the dimensionality of the data, and in our case a lot since there are many categories in each feature. This would have led to slower training and very sparse data since most of the rows would be 0 in these columns.[2]

3 Model development

3.1 Logistic Regression

Logistic regression is one of the main basic models used for binary classification. In the case involved in this project, the aim is to predict if there will be high or low availability of bikes during a specific hour, that is, aiming to predict that a given observation belongs to $Y = 1$ (Y being the target variable) based on some input features (x). Given a linear classifier: $f(x; \theta) = \text{sign}(x^T \theta)$. To train the model, the idea is to optimize the parameters (θ) by minimizing the logistic loss function (L) (another approach would be to maximize the likelihood): $L(x, y; \theta) = \ln[1 + \exp(-yx^T \theta)]$. After obtaining the probability $\hat{p} = f(x; \theta)$, the prediction is converted to a binary class label using a threshold t (commonly 0.5). The logistic regression algorithm was implemented using the sk-learn library in python. Through grid search performed with 5 cross validation folds, the optimal hyperparameters for the model were found. They are detailed below.

Hyperparameters {"C": 100, "penalty": "l1", "solver": "liblinear"}

3.2 K-Nearest Neighbors

K-Nearest Neighbors, is a non-parametric supervised learning classifier, and uses proximity to make classifications. A class label is assigned using a plurality system. To decide what class a given point should be assigned we need some way to measure the distance between that point and its k neighbors. This distance creates decision boundaries, which can be visualized using Voronoi diagrams.

$$\text{Minkowski Distance} = \left(\sum_{i=1}^n |x_i - y_i|^p \right)^{\frac{1}{p}} \quad (1)$$

Using the Minkowski distance as a generalized distance formula, we can create the two most common distance formulas used in k-nearest neighbors. When $p = 2$ we have the *euclidian* distance, and when $p = 1$, we have the *manhattan* distance [3]. The k-value determines how many neighbors are checked when the model classifies a new point. The k-nearest neighbor algorithm was implemented using the scikit-learn library in python. A grid search cross validation with 5 folds was used to find the optimal hyper-parameters, and below, we can see what

these optimal hyper-parameters are (note that the other hyperparameters that scikit-learn's KNeighborsClassifier support were left to default values).

```
Hyperparameters { "algorithm": ball_tree, "metric": manhattan,
                  "n_neighbors": 23, "weights": uniform }
```

3.3 Random Forest

Random forest is an ensemble method closely related to bootstrap aggregation (bagging). Bagging is a technique that averages over multiple low-bias high-variance models to reduce variance. There are two things that separate bagging and random forest. The first thing is that bagging can be used with any base model, while random forest assumes that the base models are classification or regression trees. The second thing is the randomness that the random forest method injects when constructing the trees. Instead of considering all features x_1, \dots, x_p as splitting variables, random forest picks $q \leq p$ random splitting variables and only considers these. This random selection of q is made at each splitting point in each of the ensemble trees B . The random selection of features will make the B trees less correlated, and thus the average over the trees predictions can result in even more variance reduction than bagging.[4] The random forest algorithm was implemented with the scikit-learn library in python. The models best hyperparameters were found by doing a grid search with $K - folds = 5$. The best parameters are shown below.

```
Hyperparameters { "bootstrap": True, "max_depth": 20,
                  "max_features": sqrt, "n_estimators": 50, "min_sample_leaf": 1,
                  "min_sample_split": 2 }
```

3.4 Boosting

Boosting is a machine learning method in which multiple simple models, like small decision trees, are combined to make a stronger model. The simple model is an algorithm that generates classifiers with an error rate slightly better than random guessing (e.g., less than 0.5 in binary classification). In contrast, a strong learner with sufficient training data can produce classifiers with a very low error probability. The algorithm could be written as detailed below.

$$F(x) = \sum_{t=1}^T \alpha_t h_t(x), \quad (2)$$

where T = Total number of simple learners, α_t = Weight (or importance) of the t -th simple learner and $h_t(x)$ = The t -th simple learner.

3.4.1 XGBoost

XGBoost (Extreme Gradient Boosting) was chosen as one of the boosting algorithms for this problem because of its ability to handle structured datasets effectively and model complex non-linear relationships between variables. The dataset

includes a mix of temporal features, such as `hour_of_day` and `day_of_week`, and weather-related variables like `temp`, `humidity`, and `rain`. XGBoost is well suited for datasets of this nature, as it captures interactions between these features without requiring extensive preprocessing.

```
Hyperparameters { "gamma": 0.1, "max_depth": 10, "reg_alpha": 1,
                  "learning_rate": 0.4, "n_estimators": 20, "reg_lambda": 0 }
```

3.4.2 CatBoost

Categorical Boosting was also considered for this problem due to its ability to handle categorical features directly. Unlike XGBoost, CatBoost processes categorical data natively and employs ordered boosting, which mitigates target leakage and enhances the robustness of the model. This is especially useful in this scenario where the dataset is quite small and overfitting can be a concern. After performing grid search over 5 cross validation folds, the optimal hyperparameters found were the ones stated below.

```
Hyperparameters { "depth": 6, "iterations": 1000,
                  "learning_rate": 0.01 }
```

4 Model Selection & Results

All the results are obtained from models trained on 80% of the dataset, and then evaluated on the remaining 20%. Random state/seed = 0 was used throughout the code, both for the dataset splits, as well as for the models that accept a random state as an argument.

Table 3: Performance metrics for all of the models considered.

Model	Accuracy	Precision	Recall	F1-Score
Logistic Regression	0.8688	0.6538	0.5862	0.6182
KNN	0.8656	0.6531	0.5517	0.5981
Random Forest	0.8812	0.6852	0.6379	0.6607
XGBoost	0.8656	0.6415	0.5862	0.6126
CatBoost	0.8781	0.7021	0.5690	0.6286
Naive Classifier	0.8187	0	0	0

Looking at the ROC Curves, and the respective areas under the curves, it can be seen that *CatBoost* performs the best. But looking at all other metrics; accuracy, precision, recall, and f1-score, as outlined in Table 3, Random Forest performs the best on all of them, except precision where CatBoost is the best. Due to the nature of the problem, recall was deemed to be important to prioritize. In this case, a higher recall value implies that the model does not wrongly predict the negative class, which is low bike demand in this case. In a business perspective, it is bad to predict low demand when in the reality, the demand is high, as

customers will chose other modes of transportation. Seeing as Random Forest displays the highest results on the most important metrics presented above, and a quite significantly higher recall than the other models, this model will be used in production.




Figure 3: ROC curves of all of the models considered.

5 Conclusions

This study demonstrates the effectiveness of machine learning models in predicting whether there will be high or low demand of bikes for Washington DC’s Capital Bikeshare system. While all the contemplated models outperform the naive classifier, Random Forest stands out from the rest by achieving the best performance in most metrics, including accuracy, precision, and F1-Score. With an AUC of 0.917, it is only slightly behind CatBoost’s top score of 0.925. However, Random Forest’s recall is 12% higher than CatBoost’s, representing a more optimal trade-off between accuracy and responsiveness to high-demand conditions. In this project recall is particularly critical, as it aligns with the goal of minimizing instances where high demand is incorrectly predicted as low, reducing the risk of bike shortages during peak demand.

In conclusion, the results evidence that Random Forest is the most reliable choice for deployment among all the models studied. This model provides a robust, data-driven solution for optimizing bike availability and promoting the use of sustainable transportation in Washington DC.

References

- [1] G. Developers, “Normalization,” <https://developers.google.com/machine-learning/crash-course/numerical-data/normalization>, 2024, accessed on December 4, 2024.
- [2] GeeksforGeeks, “One hot encoding,” <https://www.geeksforgeeks.org/ml-one-hot-encoding/>, accessed on December 4, 2024.
- [3] IBM, “What is the k-nearest neighbors (knn) algorithm?” 2024, accessed: 2024-12-08. [Online]. Available: <https://www.ibm.com/topics/knn>
- [4] A. Lindholm, N. Wahlström, F. Lindsten, and T. B. Schön, *Machine Learning - A First Course for Engineers and Scientists*. Cambridge University Press, 2022. [Online]. Available: <https://smlbook.org>

A Appendix

A.1 dataloader.py

This library was created and used for loading and splitting the dataset.

```
1 import pandas as pd
2 import numpy as np
3 from sklearn.model_selection import train_test_split
4
5 def load_and_split_data(filepath, target_column, class_zero,
6     test_size=0.2, random_state=0, cat_features=None):
7     """
8     Loads data from a CSV file, processes it, and returns train-
9     test splits.
10
11     Parameters:
12         filepath (str): Path to the CSV file.
13         target_column (str): The name of the target column to be
14         predicted.
15         class_zero (str): The name of the class to be used as the
16         reference class (0).
17         convert_cat_target (bool): Whether to convert the target
18         column to binary.
19         test_size (float): Fraction of the data to use as test set.
20         random_state (int): Random seed for reproducibility.
21         cat_features (list): List of categorical features to be
22         converted to category type.
23
24     Returns:
25         tuple: X_train, X_test, y_train, y_test
26     """
27
28     df = pd.read_csv(filepath)
29
30     # Assign 0 to the class_zero and 1 to the other class in the
31     # target column
32     df[target_column] = np.where(df[target_column] == class_zero,
33     0, 1)
34
35     if cat_features:
36         for feature in cat_features:
37             df[feature] = df[feature].astype('category')
38
39     # Split into features and target
40     X = df.copy()
41     y = X.pop(target_column)
42
43     # Split into train and test sets
44     X_train, X_test, y_train, y_test = train_test_split(X, y,
45     test_size=test_size, random_state=random_state)
46
47     return X_train, X_test, y_train, y_test
```

A.2 feature_eng.py

This library was created and used for preprocessing the dataset.

```
1 import pandas as pd
2
3
4 def preprocess(path_in, path_out, name_out):
5     """
6     Preprocess the dataset and saves it as a .csv file.
7
8     Args:
9         path_in (str): path to the dataset
10        path_out (str): path to save the preprocessed dataset
11        name_out (str): name for the preprocessed dataset
12    """
13
14    # Load the dataset
15    data = pd.read_csv(path_in)
16
17    # Define the numerical features
18    num_features = ['temp', 'dew', 'humidity', 'windspeed', '
19                    cloudcover', 'visibility']
20
21    # Drop holiday and snow columns
22    data = data.drop(columns=['holiday', 'snow'])
23
24    # Add a binary feature called "day" where 1 means "hour_of_day"
25    # is between 7 and 20, and 0 otherwise
26    data['day'] = ((data['hour_of_day'] >= 7) & (data['hour_of_day'
27    ] <= 20)).astype(int)
28
29    # Encode "snowdepth" as a binary feature where 1 means if there
30    # is snow and 0 otherwise
31    data['snowdepth'] = (data['snowdepth'] > 0).astype(int)
32
33    # Add a binary feature called "rain" where 1 means if "precip"
34    # is greater than 0, and 0 otherwise
35    data['rain'] = (data['precip'] > 0).astype(int)
36
37    # Drop "precip" column
38    data = data.drop(columns=['precip'])
39
40    # Normalize the numerical features
41    for feature in num_features:
42        data[feature] = (data[feature] - data[feature].mean()) /
43        data[feature].std()
44
45    # Save the preprocessed dataset as csv
46    data.to_csv(path_out + name_out, index=False)
```

A.3 utils.py

This library was created and used for developing and training the models.

```
1
2 from sklearn.metrics import accuracy_score, precision_score,
  recall_score, f1_score, roc_auc_score, confusion_matrix,
  roc_curve
3 from sklearn.model_selection import GridSearchCV
4 import json
5 import os
6 import matplotlib.pyplot as plt
7 import numpy as np
8 import pandas as pd
9
10
11
12
13 def find_optimal_hyperparameters(model, param_grid, X_train,
  y_train, cv=5, scoring='accuracy', n_jobs=-1, save_dir="",
  save_file='knn_best_params.json', extra_args={},
  verbose_training=False):
14
15     """
16     Find the optimal hyperparameters for a model using GridSearchCV
17
18     Parameters:
19         model: The model class
20         param_grid (dict): The hyperparameters to search over
21         X_train (pd.DataFrame): The training data
22         y_train (pd.Series): The training labels
23         cv (int): The number of cross-validation folds
24         scoring (str): The scoring metric
25         n_jobs (int): The number of jobs to run in parallel
26         save_dir (str): The directory to save the best parameters
27         save_file (str): The file to save the best parameters
28         extra_args (dict): Extra arguments to pass to the model
29         verbose_training (bool): Whether to print the training
  progress
30
31     Returns:
32         dict: The best hyperparameters found
33     """
34
35     model = model(**extra_args)
36     gs_cv = GridSearchCV(model, param_grid, cv=cv, scoring=scoring,
  n_jobs=n_jobs)
37
38     if 'CatBoostClassifier' in str(model):
39         gs_cv.fit(X_train, y_train, verbose=verbose_training)
40     else:
41         gs_cv.fit(X_train, y_train)
42
43     best_params = gs_cv.best_params_
44     print("Best parameters found: ", best_params)
45
46     if extra_args:
47         best_params.update(extra_args)
```

```

48
49     if save_dir:
50         print("Saving best parameters to '{}'.format(os.path.join(
51             save_dir, save_file).replace('\\', '/').strip()))
52         with open(os.path.join(save_dir, save_file), 'w') as f:
53             json.dump(best_params, f)
54
55     return best_params
56
57 def load_model_from_json(model, json_file):
58     """
59     Load a model from a json file
60
61     Parameters:
62         model: The model class to load
63         json_file (str): The path to the json file
64
65     Returns:
66         model: The model loaded from the json file
67     """
68
69     with open(json_file, 'r') as f:
70         params = json.load(f)
71
72     model = model(**params)
73
74     return model
75
76 def plot_roc_curves(results):
77     """
78     Plot the ROC curves for the models
79
80     Parameters:
81         results (dict): The results from fit_and_evaluate_multiple
82
83     Preconditions:
84         - results contains the fpr and tpr for each model, as well
85           as the roc_auc
86     """
87     plt.figure(figsize=(10, 7))
88     for model_name, metrics in results.items():
89         plt.plot(metrics["fpr"], metrics["tpr"], label=f"{
90             model_name} (AUC = {metrics['roc_auc']:.3f})")
91
92     plt.plot([0, 1], [0, 1], 'k--', label="Random Guess")
93     plt.title("ROC Curves")
94     plt.xlabel("False Positive Rate")
95     plt.ylabel("True Positive Rate")
96     plt.legend(loc="lower right")
97     plt.grid()
98     plt.show()
99
100 def fit_and_evaluate_multiple(models, X_train, y_train, X_test,
101     y_test, verbose=False, verbose_training=False, float_precision
102     =4):
103     """

```

```

100 Fits multiple models on the given data and evaluates them on
    the testing data.
101
102 Parameters:
103     models (list): The models
104     X_train (pd.DataFrame): The training data
105     y_train (pd.Series): The training labels
106     X_test (pd.DataFrame): The testing data
107     y_test (pd.Series): The testing labels
108     verbose (bool): Whether to print the results
109     verbose_training (bool): Whether to print the training
    progress
110     float_precision (int): The number of decimal places to
    print
111
112 Returns:
113     dict: The accuracy, precision, recall, F1, ROC AUC, and
    confusion matrix for each model
114     """
115
116 results = {}
117 for model in models:
118     results[model.__class__.__name__] = fit_and_evaluate(model,
    X_train, y_train, X_test, y_test, verbose, verbose_training,
    float_precision)
119 return results
120
121 def fit_and_evaluate(model, X_train, y_train, X_test, y_test,
    verbose=False, verbose_training=False, float_precision=4):
122     """
123     Fits a model on the given data and evaluates it on the testing
    data.
124
125     Parameters:
126         model: The model
127         X_train (pd.DataFrame): The training data
128         y_train (pd.Series): The training labels
129         X_test (pd.DataFrame): The testing data
130         y_test (pd.Series): The testing labels
131         verbose (bool): Whether to print the results
132         verbose_training (bool): Whether to print the training
    progress
133         float_precision (int): The number of decimal places to
    print
134
135     Returns:
136         dict: The accuracy, precision, recall, F1, ROC AUC,
    confusion matrix, fpr, tpr
137         """
138
139     if 'CatBoostClassifier' in str(model):
140         model.fit(X_train, y_train, verbose=verbose_training)
141     else:
142         model.fit(X_train, y_train)
143     return evaluate(model, X_test, y_test, verbose, float_precision
    )
144

```

```

145 def evaluate(model, X_test, y_test, verbose=False, float_precision
146               =4):
147     """
148     Evaluates a model on the given data and returns the accuracy,
149     precision, recall, F1, ROC AUC, and confusion matrix, fpr, tpr
150     as a dictionary.
151
152     Parameters:
153     model: The model
154     X_test (pd.DataFrame): The testing data
155     y_test (pd.Series): The testing labels
156     verbose (bool): Whether to print the results
157     float_precision (int): The number of decimal places to
158     print
159     Returns:
160     dict: The accuracy, precision, recall, F1, ROC AUC,
161     confusion matrix, fpr, tpr
162     """
163
164     y_pred = model.predict(X_test)
165     y_pred_prob = model.predict_proba(X_test)[: , 1]
166
167     acc = accuracy_score(y_test, y_pred)
168     precision = precision_score(y_test, y_pred, zero_division=0)
169     recall = recall_score(y_test, y_pred, zero_division=0)
170     f1 = f1_score(y_test, y_pred, zero_division=0)
171     roc_auc = roc_auc_score(y_test, y_pred_prob)
172     cm = confusion_matrix(y_test, y_pred)
173     fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
174
175     if verbose:
176         print(f"Evaluating {model.__class__.__name__}")
177         print(f"Accuracy: {acc:.{float_precision}f}")
178         print(f"Precision: {precision:.{float_precision}f}")
179         print(f"Recall: {recall:.{float_precision}f}")
180         print(f"F1: {f1:.{float_precision}f}")
181         print(f"ROC AUC: {roc_auc:.{float_precision}f}")
182         print(f"Confusion Matrix: \n{cm}")
183         print()
184
185     results_dict = {
186         "accuracy": acc,
187         "precision": precision,
188         "recall": recall,
189         "f1": f1,
190         "roc_auc": roc_auc,
191         "confusion_matrix": cm,
192         "fpr": fpr,
193         "tpr": tpr,
194     }
195
196     return results_dict
197
198 def fit_and_save_predictions(model, training_data, X_eval,
199                             target_column, class_zero):
200     """

```

```

196     Fits a model on the given data and saves the predictions on the
197         evaluation data.
198
199     Parameters:
200         model: The model
201         training_data (csv): Path to the training data
202         X_eval (csv): Path to the evaluation data
203         target_column (str): The name of the target column to be
204             predicted
205         class_zero (str): The name of the class to be used as the
206             reference class (0)
207     """
208
209     # Load training data
210     training_data = pd.read_csv(training_data)
211
212     # Assign 0 to the class_zero and 1 to the other class in the
213     # target column
214     training_data[target_column] = np.where(training_data[
215         target_column] == class_zero, 0, 1)
216
217     # Split training data into features and target
218     X_train = training_data.copy()
219     y_train = X_train.pop(target_column)
220
221     # Load evaluation data
222     X_eval = pd.read_csv(X_eval)
223
224     # Fit the model on the training data
225     model.fit(X_train, y_train)
226
227     # Compute the predictions on the evaluation data
228     y_pred = model.predict(X_eval)
229
230     # Reshape the predictions to a single row
231     y_pred_row = np.reshape(y_pred, (1, -1))
232
233     # Save the predictions to a CSV file
234     y_pred_df = pd.DataFrame(y_pred_row)
235     y_pred_df.to_csv("data/final_predictions.csv", header=False,
236         index=False)

```